

CS5182 Computer Graphics

NeRF & 3D Gaussian Splatting

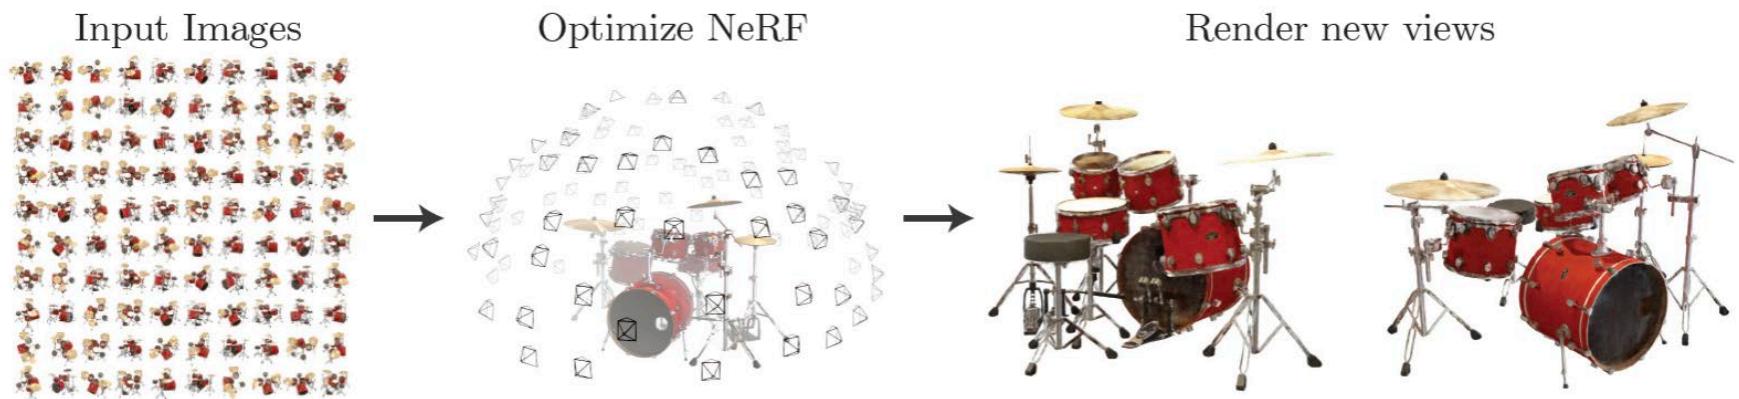
Tutorial



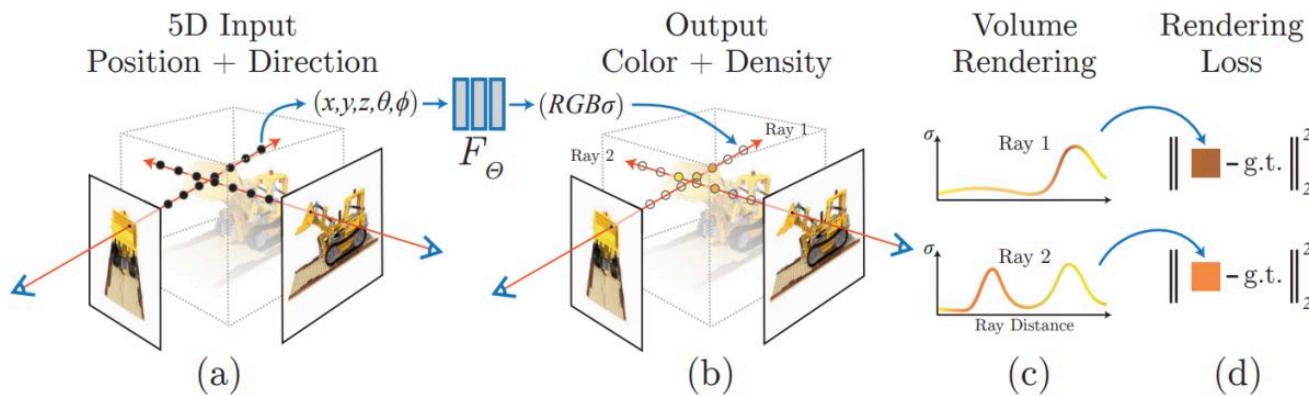
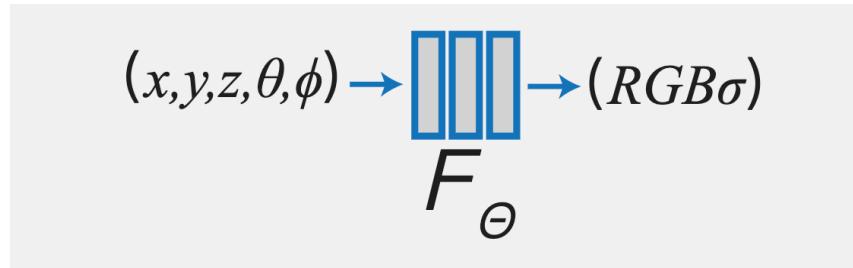
2024/25 Semester A

City University of Hong Kong (DG)

Introduction to NeRF



How does NeRF work?



Volumetric Rendering

$C(\mathbf{r})$ of camera ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$ with near and far bounds t_n and t_f is:

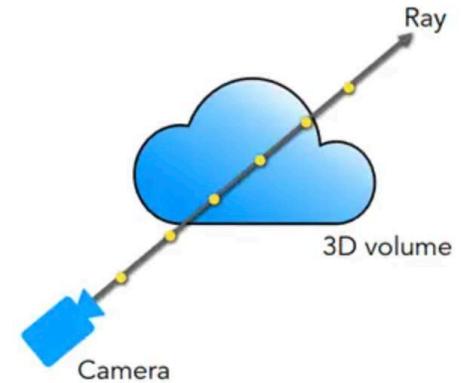
$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t)\sigma(\mathbf{r}(t))\mathbf{c}(\mathbf{r}(t), \mathbf{d})dt, \text{ where } \underline{T(t)} = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s))ds\right).$$

transmittance

-- Implementation

$$t_i \sim \mathcal{U}\left[t_n + \frac{i-1}{N}(t_f - t_n), t_n + \frac{i}{N}(t_f - t_n)\right].$$

$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i(1 - \exp(-\sigma_i \delta_i))\mathbf{c}_i, \text{ where } T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$



Optimization

-- Positional encoding for data representation

$$\gamma(p) = (\sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p)).$$

-- Hierarchical volume sampling strategy

- Coarse to fine
- Total $N_c + N_f$ samples

-- Loss function

$$\mathcal{L} = \sum_{\mathbf{r} \in \mathcal{R}} \left[\left\| \hat{C}_c(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 + \left\| \hat{C}_f(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2 \right]$$

NeRF-PyTorch

Code Issues Pull requests Actions Projects Security Insights

master 4 Branches 0 Tags Go to file Code

yenchelin Merge pull request #79 from msaroufim/patch-1 · 63a5a63 · 2 years ago · 31 Commits

config	Add configs	4 years ago
imgs	v0.1 release to public	4 years ago
.gitignore	v0.1 release to public	4 years ago
.gitmodules	v0.1 release to public	4 years ago
LICENSE	v0.1 release to public	4 years ago
README.md	Update README.md	3 years ago
download_example_data.sh	fix urls in download_example_data.sh	3 years ago

About

A PyTorch implementation of NeRF (Neural Radiance Fields) that reproduces the results.

Readme MIT license Activity 4.8k stars 55 watching 974 forks Report repository

Releases

<https://github.com/yenchelin/nerf-pytorch>

Installation

Installation

```
git clone https://github.com/yenchenlin/nerf-pytorch.git  
cd nerf-pytorch  
pip install -r requirements.txt
```



▼ Dependencies (click to expand)

Dependencies

- PyTorch 1.4
- matplotlib
- numpy
- imageio
- imageio-ffmpeg
- configargparse

The LLFF data loader requires ImageMagick.

You will also need the [LLFF code](#) (and COLMAP) set up to compute poses if you want to run on your own real data.

Codes

📁 configs	Add configs	4 years ago
📁 imgs	v0.1 release to public	
📄 .gitignore	v0.1 release to public	
📄 .gitmodules	v0.1 release to public	
📄 LICENSE	v0.1 release to public	
📄 README.md	Update README.md	
📄 download_example_data.sh	fix urls in download_example_data.sh	
📄 load_LINEMOD.py	Add load_LINEMOD.py	
📄 load_blender.py	Update load_blender.py	
📄 load_deepvoxels.py	v0.1 release to public	
📄 load_llff.py	v0.1 release to public	
📄 requirements.txt	Fix torch version to 1.11	
📄 run_nerf.py	Fix实习生问题	
📄 run_nerf_helpers.py	Update run_nerf_helpers.py	

Main running code

func batchify
func ret
func run_network
func batchify_rays
func render
func render_path
func create_nerf
func raw2outputs
func render_rays
func config_parser
func train

Codes

📁 configs	Add configs
📁 imgs	v0.1 release to public
📄 .gitignore	v0.1 release to public
📄 .gitmodules	v0.1 release to public
📄 LICENSE	v0.1 release to public
📄 README.md	Update README.md
📄 download_example_data.sh	fix urls in download_example_data.sh
📄 load_LINEMOD.py	Add load_LINEMOD.py
📄 load_blender.py	Update load_blender.py
📄 load_deevoxels.py	v0.1 release to public
📄 load_llff.py	v0.1 release to public
📄 requirements.txt	Fix torch version to 1.11
📄 run_nerf.py	Fix intrinsics problem
📄 run_nerf_helpers.py	Update run_nerf_helpers.py

class Embedder

func __init__

func create_embedding_fn

func embed

func get_embedder

class NeRF

func __init__

func forward

func load_weights_from_keras

func get_rays

func get_rays_np

func ndc_rays

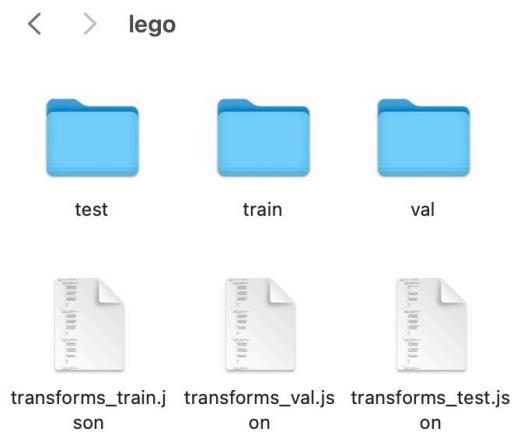
func sample_pdf

Provide the basic functions will be needed in training and rendering

Codes

📁 configs	Add configs	4 years ago
📁 imgs	v0.1 release to public	4 years ago
📄 .gitignore	v0.1 release to public	4 years ago
📄 .gitmodules	v0.1 release to public	4 years ago
📄 LICENSE	v0.1 release to public	4 years ago
📄 README.md	Update README.md	3 years ago
📄 download_example_data.sh	fix urls in download_example_data.sh	3 years ago
📄 load_LINEMOD.py	Add load_LINEMOD.py	3 years ago
📄 load_blender.py	Update load_blender.py	3 years ago
📄 load_deevoxels.py	v0.1 release to public	4 years ago
📄 load_llff.py	v0.1 release to public	4 years ago
📄 requirements.txt	Fix torch version to 1.11	2 years ago
📄 run_nerf.py	Fix intrinsics problem	3 years ago
📄 run_nerf_helpers.py	Update run_nerf_helpers.py	3 years ago

Dataloader



Dataloader

```
def load_blender_data(basedir, half_res=False, testskip=1):
    splits = ['train', 'val', 'test']
    metas = {}  Loops through each split, loading transformation metadata from JSON files corresponding to each split
    for s in splits:
        with open(os.path.join(basedir, 'transforms_{}.json'.format(s)), 'r') as fp:
            metas[s] = json.load(fp)
    all_imgs = []
    all_poses = []
    counts = [0]
    for s in splits:
        meta = metas[s]
        imgs = []
        poses = []
        if s=='train' or testskip==0:
            skip = 1
        else:
            skip = testskip
        for frame in meta['frames'][::skip]:  For each frame in the current split, reads the image and pose matrix
            fname = os.path.join(basedir, frame['file_path'] + '.png')
            imgs.append(imageio.imread(fname))
            poses.append(np.array(frame['transform_matrix']))
        imgs = (np.array(imgs) / 255.).astype(np.float32) # keep all 4 channels (RGBA)
        poses = np.array(poses).astype(np.float32)           Converts to NumPy array
        counts.append(counts[-1] + imgs.shape[0])
        all_imgs.append(imgs)
        all_poses.append(poses)
    i_split = [np.arange(counts[i], counts[i+1]) for i in range(3)]
    imgs = np.concatenate(all_imgs, 0)
    poses = np.concatenate(all_poses, 0)
    H, W = imgs[0].shape[:2]
    camera_angle_x = float(meta['camera_angle_x'])
    focal = .5 * W / np.tan(.5 * camera_angle_x)
    render_poses = torch.stack([pose_spherical(angle, -30.0, 4.0) for angle in np.linspace(-180,180,40+1)[:-1]], 0)
    if half_res:      Generates a series of camera poses for rendering, creating a spherical path around the scene
        H = H//2
        W = W//2
        focal = focal/2.
        imgs_half_res = np.zeros((imgs.shape[0], H, W, 4))
        for i, img in enumerate(imgs):
            imgs_half_res[i] = cv2.resize(img, (W, H), interpolation=cv2.INTER_AREA)
        imgs = imgs_half_res
    return imgs, poses, render_poses, [H, W, focal], i_split
```

```
"camera_angle_x": 0.6911112070083618,
"frames": [
    {
        "file_path": "./train/r_0",
        "rotation": 0.012566370614359171,
        "transform_matrix": [
            [
                -0.9999021887779236,
                0.004192245192825794,
                -0.013345719315111637,
                -0.05379832163453102
            ],
            [
                -0.013988681137561798,
                -0.2996590733528137,
                0.95394366979599,
                3.845470428466797
            ],
            [
                -4.656612873077393e-10,
                0.9540371894836426,
                0.29968830943107605,
                1.2080823183059692
            ],
            [
                0.0,
                0.0,
                0.0,
                1.0
            ]
        ],
        "i": 0
    }
]
```

Positional Encoding

```
class Embedder:
    def __init__(self, **kwargs):
        self.kwargs = kwargs
        self.create_embedding_fn()

    def create_embedding_fn(self):
        embed_fns = []
        d = self.kwargs['input_dims']
        out_dim = 0
        if self.kwargs['include_input']:
            embed_fns.append(lambda x : x)
            out_dim += d

        max_freq = self.kwargs['max_freq_log2']
        N_freqs = self.kwargs['num_freqs']

        if self.kwargs['log_sampling']:
            freq_bands = 2.**torch.linspace(0., max_freq, steps=N_freqs)
        else:
            freq_bands = torch.linspace(2.**0., 2.**max_freq, steps=N_freqs)

        for freq in freq_bands:
            for p_fn in self.kwargs['periodic_fns']:
                embed_fns.append(lambda x, p_fn=p_fn, freq=freq : p_fn(x * freq))
                out_dim += d

        self.embed_fns = embed_fns
        self.out_dim = out_dim

    def embed(self, inputs):
        return torch.cat([fn(inputs) for fn in self.embed_fns], -1)
```

Defines frequency bands based on the maximum frequency and the number of frequencies

Uses log sampling if specified, else linspace

Iterates over each frequency band and each periodic function (sine and cosine)

NeRF Module

```
class NeRF(nn.Module):
    def __init__(self, D=8, W=256, input_ch=3, input_ch_views=3, output_ch=4, skips=[4], use_viewdirs=False):
        super(NeRF, self).__init__()
        self.D = D
        self.W = W
        self.input_ch = input_ch
        self.input_ch_views = input_ch_views
        self.skips = skips
        self.use_viewdirs = use_viewdirs
        self.pts_linears = nn.ModuleList()
        [nn.Linear(input_ch, W)] + [nn.Linear(W, W) if i not in self.skips else nn.Linear(W + input_ch, W) for i in range(D-1)]
        self.views_linears = nn.ModuleList([nn.Linear(input_ch_views + W, W//2)])
        if use_viewdirs:
            self.feature_linear = nn.Linear(W, W)
            self.alpha_linear = nn.Linear(W, 1)
            self.rgb_linear = nn.Linear(W//2, 3)
        else:
            self.output_linear = nn.Linear(W, output_ch)

    def forward(self, x):
        input_pts, input_views = torch.split(x, [self.input_ch, self.input_ch_views], dim=-1)
        h = input_pts
        for i, l in enumerate(self.pts_linears):
            h = self.pts_linears[i](h)
            h = F.relu(h)
            if i in self.skips:
                h = torch.cat([input_pts, h], -1)
        if self.use_viewdirs:
            alpha = self.alpha_linear(h)
            feature = self.feature_linear(h)
            h = torch.cat([feature, input_views], -1)
            for i, l in enumerate(self.views_linears):
                h = self.views_linears[i](h)
                h = F.relu(h)
            rgb = self.rgb_linear(h)
            outputs = torch.cat([rgb, alpha], -1)
        else:
            outputs = self.output_linear(h)
        return outputs
```

Render Rays

```
def render_rays(ray_batch,
               network_fn,
               network_query_fn,
               N_samples,
               retraw=False,
               lindisp=False,
               perturb=0.,
               N_importance=0,
               network_fine=None,
               white_bkgd=False,
               raw_noise_std=0.,
               verbose=False,
               pytest=False):
    N_rays = ray_batch.shape[0] # Extracts the number of rays and their origins, directions, and bounds from the ray_batch
    rays_o, rays_d = ray_batch[:,0:3], ray_batch[:,3:6] # [N_rays, 3] each
    viewdirs = ray_batch[:, -3:] if ray_batch.shape[-1] > 8 else None
    bounds = torch.reshape(ray_batch[...,6:8], [-1,1,2])
    near, far = bounds[...,0], bounds[...,1] # [-1,1]
    t_vals = torch.linspace(0., 1., steps=N_samples)
    if not lindisp:
        z_vals = near * (1.-t_vals) + far * (t_vals) # Sample points along rays
    else:
        z_vals = 1./(1./near * (1.-t_vals) + 1./far * (t_vals))
    z_vals = z_vals.expand([N_rays, N_samples])

    if perturb > 0.:
        # get intervals between samples
        mids = .5 * (z_vals[...,-1:] + z_vals[...,:-1])
        upper = torch.cat([mids, z_vals[...,-1:]], -1)
        lower = torch.cat([z_vals[...,:1], mids], -1) # If perturbation is enabled, calculates random stratified samples between the
        # stratified samples in those intervals computed z_vals
        t_rand = torch.rand(z_vals.shape)
        # Pytest, overwrite u with numpy's fixed random numbers
        if pytest:
            np.random.seed(0)
            t_rand = np.random.rand(*list(z_vals.shape))
            t_rand = torch.Tensor(t_rand)
        z_vals = lower + (upper - lower) * t_rand
    Computes the 3D points along each ray
    pts = rays_o[...,:,None] + rays_d[...,:,None] * z_vals[...,:,None] # [N_rays, N_samples, 3]
    raw = network_query_fn(pts, viewdirs, network_fn) # Queries the neural network, processes to RGB colors, disparity, and accumulated opacity
    rgb_map, disp_map, acc_map, weights, depth_map = raw2outputs(raw, z_vals, rays_d, raw_noise_std, white_bkgd, pytest=pytest)
```

Render Rays

```
if N_importance > 0:      If using importance sampling, additional samples are generated based on the midpoint of the original z_vals and the weights computed from the first pass. These points are then queried again to refine the RGB and alpha estimates.
    rgb_map_0, disp_map_0, acc_map_0 = rgb_map, disp_map, acc_map

    z_vals_mid = .5 * (z_vals[...,1:] + z_vals[...,:-1])
    z_samples = sample_pdf(z_vals_mid, weights[...,1:-1], N_importance, det=(perturb==0.), pytest=pytest)
    z_samples = z_samples.detach()

    z_vals, _ = torch.sort(torch.cat([z_vals, z_samples], -1), -1)
    pts = rays_o[...,None,:] + rays_d[...,:,None] * z_vals[...,:,None] # [N_rays, N_samples + N_importance, 3]

    run_fn = network_fn if network_fine is None else network_fine
    raw = network_query_fn(pts, viewdirs, run_fn)

    rgb_map, disp_map, acc_map, weights, depth_map = raw2outputs(raw, z_vals, rays_d, raw_noise_std, white_bkgd, pytest=pytes

ret = {'rgb_map' : rgb_map, 'disp_map' : disp_map, 'acc_map' : acc_map}
if retraw:
    ret['raw'] = raw
if N_importance > 0:
    ret['rgb0'] = rgb_map_0           Constructs a dictionary to return the output
    ret['disp0'] = disp_map_0
    ret['acc0'] = acc_map_0
    ret['z_std'] = torch.std(z_samples, dim=-1, unbiased=False) # [N_rays]

for k in ret:
    if (torch.isnan(ret[k]).any() or torch.isinf(ret[k]).any()) and DEBUG:
        print(f"! [Numerical Error] {k} contains nan or inf.")

return ret
```

Train

Configuration & Data Loading

```
def train():

    parser = config_parser()
    args = parser.parse_args()

    K = None

    if args.dataset_type == 'blender':      Depending on the specified dataset type, load data
        images, poses, render_poses, hwf, i_split = load_blender_data(args.datadir, args.half_res, args.testskip)
        print('Loaded blender', images.shape, render_poses.shape, hwf, args.datadir)
        i_train, i_val, i_test = i_split

        near = 2.
        far = 6.

        if args.white_bkgd:
            images = images[...,:3]*images[...,-1:] + (1.-images[...,-1:])
        else:
            images = images[...,:3]

    # Cast intrinsics to right types
    H, W, focal = hwf
    H, W = int(H), int(W)
    hwf = [H, W, focal]

    if K is None:
        K = np.array([
            [focal, 0, 0.5*W],
            [0, focal, 0.5*H],
            [0, 0, 1]
        ])
```

Train

Model Initialization

```
# Create log dir and copy the config file
basedir = args.basedir
expname = args.expname
os.makedirs(os.path.join(basedir, expname), exist_ok=True)
f = os.path.join(basedir, expname, 'args.txt')
with open(f, 'w') as file:
    for arg in sorted(vars(args)):    Creates a directory for saving training logs and configurations
        attr = getattr(args, arg)
        file.write('{} = {}\n'.format(arg, attr))
if args.config is not None:
    f = os.path.join(basedir, expname, 'config.txt')
    with open(f, 'w') as file:
        file.write(open(args.config, 'r').read())

# Create nerf model
render_kwargs_train, render_kwargs_test, start, grad_vars, optimizer = create_nerf(args)
global_step = start                                Initializes the NeRF model and optimizer
                                                    Sets the starting point for training
bds_dict = {
    'near' : near,
    'far' : far,
}
render_kwargs_train.update(bds_dict)
render_kwargs_test.update(bds_dict)
```

Train

Training Loop Preparation

```
# Prepare raybatch tensor if batching random rays
N_rand = args.N_rand
use_batching = not args.no_batching
if use_batching:
    # For random ray batching      Prepares batches of rays for training
    print('get rays')
    rays = np.stack([get_rays_np(H, W, K, p) for p in poses[:, :, :3, :4]], 0) # [N, ro+rd, H, W, 3]
    print('done, concats')
    rays_rgb = np.concatenate([rays, images[:, None]], 1) # [N, ro+rd+rgb, H, W, 3]
    rays_rgb = np.transpose(rays_rgb, [0, 2, 3, 1, 4]) # [N, H, W, ro+rd+rgb, 3]
    rays_rgb = np.stack([rays_rgb[i] for i in i_train], 0) # train images only
    rays_rgb = np.reshape(rays_rgb, [-1, 3, 3]) # [(N-1)*H*W, ro+rd+rgb, 3]
    rays_rgb = rays_rgb.astype(np.float32)
    print('shuffle rays')
    np.random.shuffle(rays_rgb)

    print('done')
    i_batch = 0

# Move testing data to GPU
render_poses = torch.Tensor(render_poses).to(device)
# Move training data to GPU
if use_batching:                                Moves data to the GPU
    images = torch.Tensor(images).to(device)
poses = torch.Tensor(poses).to(device)
if use_batching:
    rays_rgb = torch.Tensor(rays_rgb).to(device)
```

Train

Training Process:

- Ray Sampling

```
start = start + 1
for i in range(start, N_iters):
    time0 = time.time()

    # Sample random ray batch
    if use_batching:
        # Random over all images      Sampling rays from the dataset
        batch = rays_rgb[i_batch:i_batch+N_rand] # [B, 2+1, 3*?]
        batch = torch.transpose(batch, 0, 1)
        batch_rays, target_s = batch[:2], batch[2]

        i_batch += N_rand
        if i_batch >= rays_rgb.shape[0]:
            print("Shuffle data after an epoch!")
            rand_idx = torch.randperm(rays_rgb.shape[0])
            rays_rgb = rays_rgb[rand_idx]
            i_batch = 0
```

Train

Training Process:

- Volume rendering
- Optimization

```
##### Core optimization loop #####
rgb, disp, acc, extras = render(H, W, K, chunk=args.chunk, rays=batch_rays,
| | | | | Rendering RGB values and densities using the model
| | | | | verbose=i < 10, retraw=True,
| | | | | **render_kwargs_train)

optimizer.zero_grad()
img_loss = img2mse(rgb, target_s)
trans = extras['raw'][..., -1]
loss = img_loss
psnr = mse2psnr(img_loss)

if 'rgb0' in extras:
    img_loss0 = img2mse(extras['rgb0'], target_s)
    loss = loss + img_loss0
    psnr0 = mse2psnr(img_loss0)  Calculating the loss

loss.backward()
optimizer.step()  Performs backpropagation to update model weights
```

Train

Training Process:

- Learning Rate Adjustment
- Logging and Checkpoints

```
# NOTE: IMPORTANT!
### update learning rate ####
decay_rate = 0.1
decay_steps = args.lrate_decay * 1000      Optionally adjusts the learning rate based on the current
new_lrate = args.lrate * (decay_rate ** (global_step / decay_steps))
for param_group in optimizer.param_groups:
    param_group['lr'] = new_lrate
#####
dt = time.time() - time0
# print(f"Step: {global_step}, Loss: {loss}, Time: {dt}")
#####           end           #####
# Rest is logging
if i % args.i_weights == 0:
    path = os.path.join(basedir, expname, '{:06d}.tar'.format(i))
    torch.save({
        'global_step': global_step,
        'network_fn_state_dict': render_kwargs_train['network_fn'].state_dict(),
        'network_fine_state_dict': render_kwargs_train['network_fine'].state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
    }, path)                                Logs the training progress, saves model checkpoints
    print('Saved checkpoints at', path)
```

Train

Evaluation & Rendering

```
if i%args.i_video==0 and i > 0:  
    # Turn on testing mode  
    with torch.no_grad():  
        rgbs, disps = render_path(render_poses, hwf, K, args.chunk, render_kwargs_test)  
        print('Done, saving', rgbs.shape, disps.shape)  
        moviebase = os.path.join(basedir, expname, '{}_spiral_{:06d}'.format(expname, i))  
        imageio.mimwrite(moviebase + 'rgb.mp4', to8b(rgbs), fps=30, quality=8)  
        imageio.mimwrite(moviebase + 'disp.mp4', to8b(disps / np.max(disps)), fps=30, quality=8)  
  
if i%args.i_testset==0 and i > 0:  
    testsavedir = os.path.join(basedir, expname, 'testset_{:06d}'.format(i))  
    os.makedirs(testsavedir, exist_ok=True)  
    print('test poses shape', poses[i_test].shape)  
    with torch.no_grad():  
        render_path(torch.Tensor(poses[i_test]).to(device), hwf, K, args.chunk, render_kwargs_test, gt_imgs=images[i_test], savedir=testsavedir)  
    print('Saved test set')  
  
Renders images or videos from the test set or predefined poses  
  
if i%args.i_print==0:  
    tqdm.write(f"[TRAIN] Iter: {i} Loss: {loss.item()} PSNR: {psnr.item()}")  
....
```

Run

How To Run?

Quick Start

Download data for two example datasets: `lego` and `fern`

```
bash download_example_data.sh
```



To train a low-res `lego` NeRF:

```
python run_nerf.py --config configs/lego.txt
```



After training for 100k iterations (~4 hours on a single 2080 Ti), you can find the following video at

`logs/lego_test/lego_test_spiral_100000_rgb.mp4`.



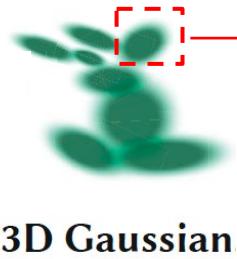
Introduction to 3D Gaussian Splatting

- Reconstructing the 3D world from images



Method

Gaussians are defined by a full 3D covariance matrix Σ defined in world space centered at point (mean) μ and is multiplied by α in blending process



- **Position (Mean μ):** location (XYZ)
- **Covariance Matrix (Σ):** rotation and scaling $\Sigma = RSS^T R^T$
- **Opacity (α):** Transparency
- **Color (RGB) or Spherical Harmonics (SH) coefficients**

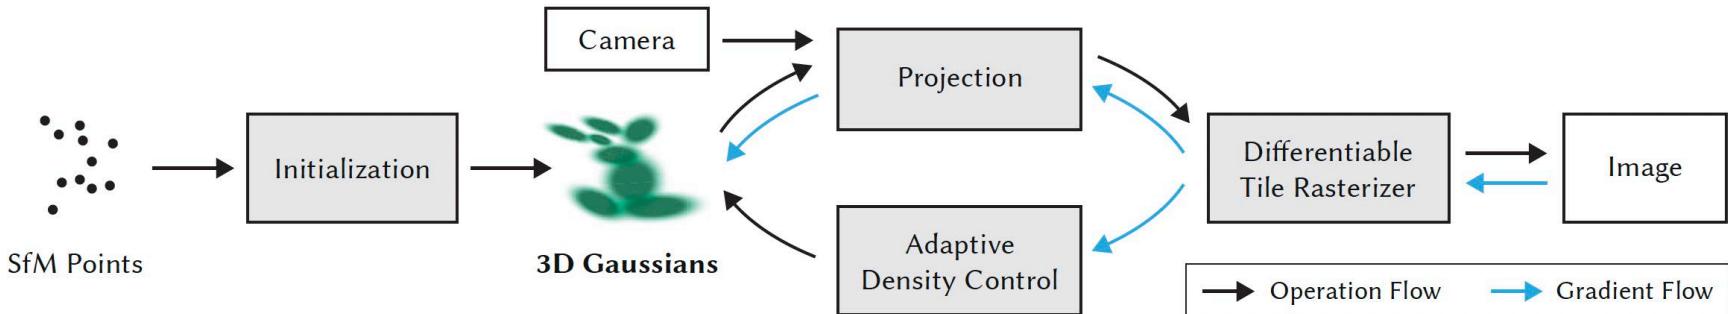
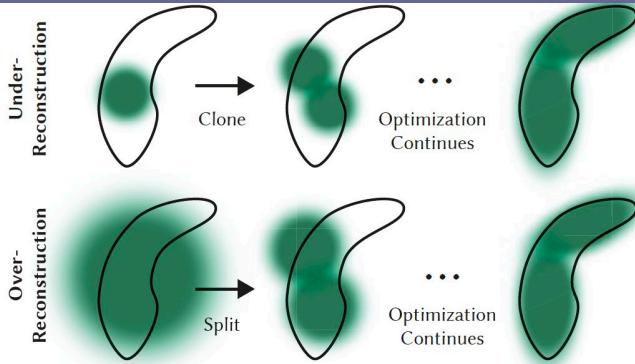


Fig. 2. Optimization starts with the sparse SfM point cloud and creates a set of 3D Gaussians. We then optimize and adaptively control the density of this set of Gaussians. During optimization we use our fast tile-based renderer, allowing competitive training times compared to SOTA fast radiance field methods. Once trained, our renderer allows real-time navigation for a wide variety of scenes.

Optimization



Adaptive Gaussian demsification

Fig. 4. Our adaptive Gaussian densification scheme. *Top row (under-reconstruction)*: When small-scale geometry (black outline) is insufficiently covered, we clone the respective Gaussian. *Bottom row (over-reconstruction)*: If small-scale geometry is represented by one large splat, we split it in two.

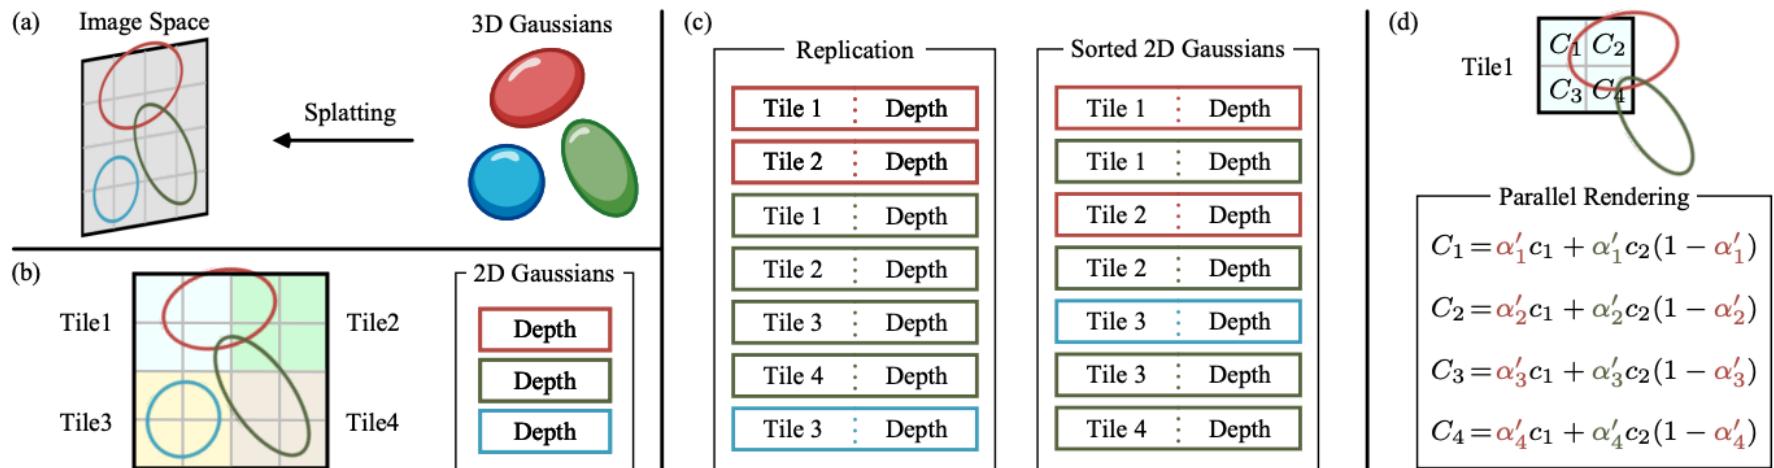


Fig. 4. An illustration of the forward process of 3D GS (see Sec. 3.1). (a) The splatting step projects 3D Gaussians into image space. (b) 3D GS divides the image into multiple non-overlapping patches, i.e., tiles. (c) 3D GS replicates the Gaussians which cover several tiles, assigning each copy an identifier, i.e., a tile ID. (d) By rendering the sorted Gaussians, we can obtain all pixels within the tile. Note that the computational workflows for pixels and tiles are independent and can be done in parallel. Best viewed in color.

Code

graphdeco-inria / gaussian-splatting + ▾ ○ !! ✉

Code Issues 427 Pull requests 28 Actions Projects Security Insights

 gaussian-splatting Public Watch 116 Fork 1.7k Star 13.6k

main 2 Branches 0 Tags t Add file Code

 alanvinx	revert rasterizer version	8a70a8c · last week	240 Commits
 SIBR_viewers @ 212f9fb	new release	3 weeks ago	
 arguments	Revert new release, this reverts commit 2130164.	last week	
 assets	Revert new release, this reverts commit 2130164.	last week	
 gaussian_renderer	Revert new release, this reverts commit 2130164.	last week	
 lapisPyTorch	Initial commit	last year	
 scene	Revert new release, this reverts commit 2130164.	last week	
 submodules	revert rasterizer version	last week	
 utils	Revert new release, this reverts commit 2130164.	last week	
 .gitignore	Initial commit	last year	
 .gitmodules	Revert new release, this reverts commit 2130164.	last week	
 LICENSE.md	License correction	7 months ago	

About

Original reference implementation of "3D Gaussian Splatting for Real-Time Radiance Field Rendering"

 [repo-sam.inria.fr/fungraph/3d-gaussian-splatting](#)

computer-vision computer-graphics radiance-field

Readme View license Activity Custom properties 13.6k stars 116 watching 1.7k forks Report repository

Releases

No releases published

<https://github.com/graphdeco-inria/gaussian-splatting>

Setup

Setup

Local Setup

Our default, provided install method is based on Conda package and environment management:

```
SET DISTUTILS_USE_SDK=1 # Windows only  
conda env create --file environment.yml  
conda activate gaussian_splatting
```



Please note that this process assumes that you have CUDA SDK **11** installed, not **12**. For modifications, see below.

Tip: Downloading packages and creating a new environment with Conda can require a significant amount of disk space. By default, Conda will use the main system hard drive. You can avoid this by specifying a different package download location and an environment on a different drive:

```
conda config --add pkgs_dirs <Drive>/<pkg_path>  
conda env create --file environment.yml --prefix <Drive>/<env_path>/gaussian_splatting  
conda activate <Drive>/<env_path>/gaussian_splatting
```



Setup (install COLMAP)

```
!sudo apt-get install \
    git \
    cmake \
    ninja-build \
    build-essential \
    libboost-program-options-dev \
    libboost-filesystem-dev \
    libboost-graph-dev \
    libboost-system-dev \
    libeigen3-dev \
    libflann-dev \
    libfreeimage-dev \
    libmetis-dev \
    libgoogle-glog-dev \
    libgtest-dev \
    libsqlite3-dev \
    libglew-dev \
    qtbase5-dev \
    libqt5opengl5-dev \
    libcgal-dev \
    libceres-dev

!git clone https://github.com/colmap/colmap.git
cd colmap
!git checkout dev
mkdir build
cd build
!cmake .. -GNinja -DCMAKE_CUDA_ARCHITECTURES=native
!ninja
!sudo ninja install
```

Images + Camera pose



3D point cloud

GaussianModel

Initialization

```
def __init__(self, sh_degree : int):
    self.active_sh_degree = 0
    self.max_sh_degree = sh_degree
    self._xyz = torch.empty(0)
    self._features_dc = torch.empty(0)
    self._features_rest = torch.empty(0)
    self._scaling = torch.empty(0)
    self._rotation = torch.empty(0)
    self._opacity = torch.empty(0)
    self.max_radii2D = torch.empty(0)
    self.xyz_gradient_accum = torch.empty(0)
    self.denom = torch.empty(0)
    self.optimizer = None
    self.percent_dense = 0
    self.spatial_lr_scale = 0
    self.setup_functions()
```

GaussianModel

Create gaussians from point cloud

```
def create_from_pcd(self, pcd : BasicPointCloud, spatial_lr_scale : float):
    self.spatial_lr_scale = spatial_lr_scale
    fused_point_cloud = torch.tensor(np.asarray(pcd.points)).float().cuda()      Converts the point cloud colors to spherical
    fused_color = RGB2SH(torch.tensor(np.asarray(pcd.colors)).float().cuda())      harmonics representation
    features = torch.zeros((fused_color.shape[0], 3, (self.max_sh_degree + 1) ** 2)).float().cuda()
    features[:, :, 0] = fused_color           Store the features of the Gaussian model
    features[:, :, 1:] = 0.0

    print("Number of points at initialisation : ", fused_point_cloud.shape[0])

    dist2 = torch.clamp_min(distCUDA2(torch.from_numpy(np.asarray(pcd.points)).float().cuda()), 0.0000001)      Computes the squared distances between points in the point cloud
    scales = torch.log(torch.sqrt(dist2))[:, None].repeat(1, 3)   Calculates the scaling factors based on the distances
    rots = torch.zeros((fused_point_cloud.shape[0], 4), device="cuda")
    rots[:, 0] = 1           Initializes rotation parameters

    opacities = inverse_sigmoid(0.1 * torch.ones((fused_point_cloud.shape[0], 1), dtype=torch.float, device="cuda"))
    Initialize opacities for the Gaussians

    self._xyz = nn.Parameter(fused_point_cloud.requires_grad_(True))
    self._features_dc = nn.Parameter(features[:, :, 0:1].transpose(1, 2).contiguous().requires_grad_(True))
    self._features_rest = nn.Parameter(features[:, :, 1:].transpose(1, 2).contiguous().requires_grad_(True))
    self._scaling = nn.Parameter(scales.requires_grad_(True))
    self._rotation = nn.Parameter(rots.requires_grad_(True))           Initializes all the necessary parameters
    self._opacity = nn.Parameter(opacities.requires_grad_(True))
    self.max_radii2D = torch.zeros((self.get_xyz.shape[0]), device="cuda")
```

GaussianModel

Adaptive Gaussian densification -- split

```
def densify_and_split(self, grads, grad_threshold, scene_extent, N=2):
    n_init_points = self.get_xyz.shape[0]
    # Extract points that satisfy the gradient condition
    padded_grad = torch.zeros((n_init_points), device="cuda")      Initializes a tensor to hold the gradient values
    padded_grad[:grads.shape[0]] = grads.squeeze()
    selected_pts_mask = torch.where(padded_grad >= grad_threshold, True, False)  Creates a mask to identify points that have
    selected_pts_mask = torch.logical_and(selected_pts_mask,
                                         torch.max(self.get_scaling, dim=1).values > self.percent_dense*scene_extent)
    Further refines the mask to include only those points whose maximum scaling value exceeds a certain density threshold relative to the scene extent

    stds = self.get_scaling[selected_pts_mask].repeat(N,1)
    means = torch.zeros((stds.size(0), 3),device="cuda")      Generates new sample points
    samples = torch.normal(mean=means, std=stds)
    rots = build_rotation(self._rotation[selected_pts_mask]).repeat(N,1,1)
    new_xyz = torch.bmm(rots, samples.unsqueeze(-1)).squeeze(-1) + self.get_xyz[selected_pts_mask].repeat(N, 1)
    new_scaling = self.scaling_inverse_activation(self.get_scaling[selected_pts_mask].repeat(N,1) / (0.8*N))
    new_rotation = self._rotation[selected_pts_mask].repeat(N,1)
    new_features_dc = self._features_dc[selected_pts_mask].repeat(N,1,1)      Updates the properties for the new Gaussian points
    new_features_rest = self._features_rest[selected_pts_mask].repeat(N,1,1)
    new_opacity = self._opacity[selected_pts_mask].repeat(N,1)

    self.densification_postfix(new_xyz, new_features_dc, new_features_rest, new_opacity, new_scaling, new_rotation)

    prune_filter = torch.cat((selected_pts_mask, torch.zeros(N * selected_pts_mask.sum(), device="cuda", dtype=bool)))
    self.prune_points(prune_filter)  Creates a filter combining the selected points mask and a zero tensor for the new points that should
                                    not be included in pruning. Calls prune_points to remove points
```

GaussianModel

Adaptive Gaussian densification – clone

```
def densify_and_clone(self, grads, grad_threshold, scene_extent):
    # Extract points that satisfy the gradient condition
    selected_pts_mask = torch.where(torch.norm(grads, dim=-1) >= grad_threshold, True, False)
    selected_pts_mask = torch.logical_and(selected_pts_mask,
                                         torch.max(self.get_scaling, dim=1).values <= self.percent_dense*scene_extent)
    Further refines the mask to include only those points whose maximum scaling value does not exceed a certain density threshold relative to the scene extent
    new_xyz = self._xyz[selected_pts_mask]
    new_features_dc = self._features_dc[selected_pts_mask]
    new_features_rest = self._features_rest[selected_pts_mask]
    new_opacities = self._opacity[selected_pts_mask]           Clone selected points
    new_scaling = self._scaling[selected_pts_mask]
    new_rotation = self._rotation[selected_pts_mask]

    self.densification_postfix(new_xyz, new_features_dc, new_features_rest, new_opacities, new_scaling, new_rotation)
```

GaussianModel

Adaptive Gaussian densification – prune

```
def densify_and_prune(self, max_grad, min_opacity, extent, max_screen_size):
    grads = self.xyz_gradient_accum / self.denom           Calculates the gradients
    grads[grads.isnan()] = 0.0

    self.densify_and_clone(grads, max_grad, extent)        Calls the two functions
    self.densify_and_split(grads, max_grad, extent)

    prune_mask = (self.get_opacity < min_opacity).squeeze()  Initializes a mask for pruning points that have an opacity below the specified minimum opacity threshold
    if max_screen_size:                                     Identifies points that have a radius greater than the maximum screen size
        big_points_vs = self.max_radii2D > max_screen_size  Identifies points whose maximum scaling value exceeds 10% of the scene extent
        big_points_ws = self.get_scaling.max(dim=1).values > 0.1 * extent
        prune_mask = torch.logical_or(torch.logical_or(prune_mask, big_points_vs), big_points_ws)
    self.prune_points(prune_mask)

    torch.cuda.empty_cache()
```

Training

```
def training(dataset, opt, pipe, testing_iterations, saving_iterations, checkpoint_iterations, checkpoint, debug_from):
    first_iter = 0      Initializes counters
    tb_writer = prepare_output_and_logger(dataset)  Prepares logging
    gaussians = GaussianModel(dataset.sh_degree)    Creates a Gaussian model
    scene = Scene(dataset, gaussians)   Creates a scene using the dataset
    gaussians.training_setup(opt)     Sets up the Gaussian model for training
    if checkpoint:
        (model_params, first_iter) = torch.load(checkpoint)  If a checkpoint is provided, it loads the model parameters and resumes from the specified
        gaussians.restore(model_params, opt)                  iteration

    bg_color = [1, 1, 1] if dataset.white_background else [0, 0, 0]
    background = torch.tensor(bg_color, dtype=torch.float32, device="cuda")  Defines the background color based on the dataset

    iter_start = torch.cuda.Event(enable_timing = True)
    iter_end = torch.cuda.Event(enable_timing = True)      Initializes CUDA events to track time
```

Training

```
viewpoint_stack = None
ema_loss_for_log = 0.0
progress_bar = tqdm(range(first_iter, opt.iterations), desc="Training progress")      Sets up a progress bar for monitoring training
first_iter += 1
for iteration in range(first_iter, opt.iterations + 1):  Iterates through the specified number of training iterations
    if network_gui.conn == None:      Attempts to connect to a GUI for real-time monitoring
        network_gui.try_connect()
    while network_gui.conn != None:
        try:
            net_image_bytes = None
            custom_cam, do_training, pipe.convert_SHs_python, pipe.compute_cov3D_python, keep_alive, scaling_modifer = network_gui.receive()
            if custom_cam != None:
                net_image = render(custom_cam, gaussians, pipe, background, scaling_modifer)["render"]
                net_image_bytes = memoryview((torch.clamp(net_image, min=0, max=1.0) * 255).byte().permute(1, 2, 0).contiguous().cpu().numpy())
            network_gui.send(net_image_bytes, dataset.source_path)
            if do_training and ((iteration < int(opt.iterations)) or not keep_alive):
                break
        except Exception as e:
            network_gui.conn = None

iter_start.record()

gaussians.update_learning_rate(iteration)

# Every 1000 its we increase the levels of SH up to a maximum degree
if iteration % 1000 == 0:
    gaussians.oneupSHdegree()
# Pick a random Camera
if not viewpoint_stack:
    viewpoint_stack = scene.getTrainCameras().copy()
viewpoint_cam = viewpoint_stack.pop(randint(0, len(viewpoint_stack)-1))      Randomly selects a camera from the available viewpoints for rendering
# Render
if (iteration - 1) == debug_from:
    pipe.debug = True
bg = torch.rand(3, device="cuda") if opt.random_background else background      Render
render_pkg = render(viewpoint_cam, gaussians, pipe, bg)
image, viewspace_point_tensor, visibility_filter, radii = render_pkg["render"], render_pkg["viewspace_points"], render_pkg["visibility_filter"], render_pkg["radii"]
# Loss
gt_image = viewpoint_cam.original_image.cuda()
Ll1 = l1_loss(image, gt_image)
loss = (1.0 - opt.lambda_dssim) * Ll1 + opt.lambda_dssim * (1.0 - ssim(image, gt_image))      Computes the loss using L1 loss and SSIM
loss.backward()
iter_end.record()
```

Training

```
with torch.no_grad():
    # Progress bar
    ema_loss_for_log = 0.4 * loss.item() + 0.6 * ema_loss_for_log
    if iteration % 10 == 0:
        progress_bar.set_postfix({"Loss": f"{ema_loss_for_log:.{7}f}"})      Perform progress bar
        progress_bar.update(10)
    if iteration == opt.iterations:
        progress_bar.close()
    # Log and save
    training_report(tb_writer, iteration, L1, loss, l1_loss, iter_start.elapsed_time(iter_end), testing_iterations, scene, render, (pipe, background))
    if (iteration in saving_iterations):
        print("\n[ITER {}] Saving Gaussians".format(iteration))          Saves the model state at specified iterations
        scene.save(iteration)
    # Densification
    if iteration < opt.densify_until_iter:                                Densification for pruning
        # Keep track of max radii in image-space for pruning
        gaussians.max_radii2D[visibility_filter] = torch.max(gaussians.max_radii2D[visibility_filter], radii[visibility_filter])
        gaussians.add_densification_stats(viewspace_point_tensor, visibility_filter)

        if iteration > opt.densify_from_iter and iteration % opt.densification_interval == 0:
            size_threshold = 20 if iteration > opt.opacity_reset_interval else None
            gaussians.densify_and_prune(opt.densify_grad_threshold, 0.005, scene.cameras_extent, size_threshold)

        if iteration % opt.opacity_reset_interval == 0 or (dataset.white_background and iteration == opt.densify_from_iter):
            gaussians.reset_opacity()

    # Optimizer step
    if iteration < opt.iterations:
        gaussians.optimizer.step()      Performs an optimization step to update the model weights
        gaussians.optimizer.zero_grad(set_to_none = True)

    if (iteration in checkpoint_iterations):
        print("\n[ITER {}] Saving Checkpoint".format(iteration))          Saves a checkpoint of the model at specified iterations
        torch.save((gaussians.capture(), iteration), scene.model_path + "/chkpnt" + str(iteration) + ".pth")
```

Render

```
def render(viewpoint_camera, pc : GaussianModel, pipe, bg_color : torch.Tensor, scaling_modifier = 1.0, override_color = None):
    """
    Render the scene.

    Background tensor (bg_color) must be on GPU!
    """

    # Create zero tensor. We will use it to make pytorch return gradients of the 2D (screen-space) means
    screenspace_points = torch.zeros_like(pc.get_xyz, dtype=pc.get_xyz.dtype, requires_grad=True, device="cuda") + 0
    try:
        screenspace_points.retain_grad()
    except:
        pass
    # Set up rasterization configuration
    tanfov_x = math.tan(viewpoint_camera.FoVx * 0.5)  Calculates the tangent of half the horizontal and vertical field of view angles to set up the rasterization setting
    tanfov_y = math.tan(viewpoint_camera.FoVy * 0.5)

    raster_settings = GaussianRasterizationSettings(
        image_height=int(viewpoint_camera.image_height),
        image_width=int(viewpoint_camera.image_width),  Initializes the rasterization settings including image dimensions, field of view,
        tanfov_x=tanfov_x,
        tanfov_y=tanfov_y,
        bg=bg_color,
        scale_modifier=scaling_modifier,
        viewmatrix=viewpoint_camera.world_view_transform,
        projmatrix=viewpoint_camera.full_proj_transform,
        sh_degree=pc.active_sh_degree,
        campos=viewpoint_camera.camera_center,
        prefiltered=False,
        debug=pipe.debug
    )
    rasterizer = GaussianRasterizer(raster_settings=raster_settings)  Instantiates a GaussianRasterizer using the defined settings

    means3D = pc.get_xyz
    means2D = screenspace_points      Extracts the 3D means (positions), 2D means (screen-space points), and opacity information from the Gaussian model
    opacity = pc.get_opacity
```

Render

```
# If precomputed 3d covariance is provided, use it. If not, then it will be computed from
# scaling / rotation by the rasterizer.
scales = None
rotations = None
cov3D_precomp = None
if pipe.compute_cov3D_python:
    cov3D_precomp = pc.get_covariance(scaling_modifier)
else:
    scales = pc.get_scaling           Checks if covariance data is precomputed or needs to be calculated during rendering
    rotations = pc.get_rotation

# If precomputed colors are provided, use them. Otherwise, if it is desired to precompute colors
# from SHs in Python, do it. If not, then SH -> RGB conversion will be done by rasterizer.
shs = None
colors_precomp = None
if override_color is None:          Determines how to handle colors: either using precomputed colors or calculating
                                    them from spherical harmonics
    if pipe.convert_SHs_python:
        shs_view = pc.get_features.transpose(1, 2).view(-1, 3, (pc.max_sh_degree+1)**2)
        dir_pp = (pc.get_xyz - viewpoint_camera.camera_center.repeat(pc.get_features.shape[0], 1))
        dir_pp_normalized = dir_pp/dir_pp.norm(dim=1, keepdim=True)
        sh2rgb = eval_sh(pc.active_sh_degree, shs_view, dir_pp_normalized)
        colors_precomp = torch.clamp_min(sh2rgb + 0.5, 0.0)
    else:
        shs = pc.get_features
else:
    colors_precomp = override_color
# Rasterize visible Gaussians to image, obtain their radii (on screen).
rendered_image, radii = rasterizer(
    means3D = means3D,
    means2D = means2D,           Calls the rasterizer to generate the rendered image and retrieve the radii of the visible Gaussians
    shs = shs,
    colors_precomp = colors_precomp,
    opacities = opacity,
    scales = scales,
    rotations = rotations,
    cov3D_precomp = cov3D_precomp)
# Those Gaussians that were frustum culled or had a radius of 0 were not visible.
# They will be excluded from value updates used in the splitting criteria.
return {"render": rendered_image,
        "viewspace_points": screenspace_points,
        "visibility_filter" : radii > 0,           Returns a dictionary containing the rendered image, the screen-space points, a
                                                visibility filter indicating which Gaussians were visible, and their radii
        "radii": radii}
```

Running & Evaluation

Running

To run the optimizer, simply use

```
python train.py -s <path to COLMAP or NeRF Synthetic dataset>
```



Evaluation

By default, the trained models use all available images in the dataset. To train them while withholding a test set for evaluation, use the `--eval` flag. This way, you can render training/test sets and produce error metrics as follows:

```
python train.py -s <path to COLMAP or NeRF Synthetic dataset> --eval # Train with train/test split  
python render.py -m <path to trained model> # Generate renderings  
python metrics.py -m <path to trained model> # Compute error metrics on renderings
```