

# CS5222 Computer Networks and Internets

## Tutorial 10 (Week 11)

Prof Weifa Liang

Weifa.liang@cityu-dg.edu.cn

Slides based on book *Computer Networking: A Top-Down Approach.*

Revision

# Distance vector algorithm

## key idea:

- from time-to-time, each node sends its own distance vector estimate to neighbors
- when  $x$  receives new DV estimate from **any neighbor**, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c_{x,v} + D_v(y)\} \text{ for each node } y \in N$$

- under natural conditions, the estimate  $D_x(y)$  converges to the actual least cost  $d_x(y)$

# Distance vector algorithm

Based on *Bellman-Ford* (BF) equation:

Bellman-Ford equation

Let  $D_x(y)$ : cost of least-cost path from  $x$  to  $y$ .

Then:


$$D_x(y) = \min_v \{ c_{x,v} + D_v(y) \}$$

$v$ 's estimated least-cost-path cost to  $y$

cost of link from  $x$  to  $v$

$\min$  taken over all neighbors  $v$  of  $x$

1. Consider the distance vector algorithm and show the distance vector entries at node z at each step of the algorithm.



Initialization:

	u	v	x	y	z
$D_u$ :	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
$D_v$ :	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
$D_x$ :	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
$D_y$ :	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$
$D_z$ :	$\infty$	5	2	10	0



	u	v	x	y	z
u	0	1	2	$\infty$	$\infty$
v	1	0	$\infty$	15	5
x	2	$\infty$	0	1	2
y	$\infty$	15	1	0	10
z	4	5	2	3	0





	u	v	x	y	z
u	0	1	2	3	4
v	1	0	3	15	5
x	2	3	0	1	2
y	3	15	1	0	3
z	4	5	2	3	0



	u	v	x	y	z
u	0	1	2	3	4
v	1	0	3	4	5
x	2	3	0	1	2
y	3	4	1	0	3
z	4	5	2	3	0

2. Consider the distance vector algorithm for the network below and show the distance vector entries of each node at each step of the algorithm.



Node x

	cost to		
	x	y	z
D <sub>x</sub>	0	2	7
D <sub>y</sub>	∞	∞	∞
D <sub>z</sub>	∞	∞	∞

Node y

	cost to		
	x	y	z
D <sub>x</sub>	∞	∞	∞
D <sub>y</sub>	2	0	1
D <sub>z</sub>	∞	∞	∞

Node z

	cost to		
	x	y	z
D <sub>x</sub>	∞	∞	∞
D <sub>y</sub>	∞	∞	∞
D <sub>z</sub>	7	1	0

	cost to		
	x	y	z
from x	0	2	3
from y	2	0	1
from z	7	1	0

	cost to		
	x	y	z
from x	0	2	7
from y	2	0	1
from z	7	1	0


	cost to		
	x	y	z
from x	0	2	7
from y	2	0	1
from z	3	1	0

	cost to		
	x	y	z
from x	0	2	3
from y	2	0	1
from z	3	1	0


	cost to		
	x	y	z
from x	0	2	3
from y	2	0	1
from z	3	1	0

	cost to		
	x	y	z
from x	0	2	3
from y	2	0	1
from z	3	1	0

time



3. Consider the network fragment shown below. x has only two attached neighbors, w and y. w has a minimum-cost path to destination u (not shown) of 5, and y has a minimum-cost path to u of 6. The complete paths from w and y to u (and between w and y) are not shown. All link costs in the network have strictly positive integer values. Give x's distance vector for destinations w, y, and u.



Answer:

Given  $D_w(u) = 5, D_y(u) = 6,$


$$D_x(w) = \min \{c_{x,w} + D_w(w), c_{x,y} + D_y(w)\} = \min\{2 + 0, 5 + 2\} = 2$$

$$D_x(y) = \min \{c_{x,w} + D_w(y), c_{x,y} + D_y(y)\} = \min\{2 + 2, 5 + 0\} = 4$$

$$D_x(u) = \min\{c_{x,w} + D_w(u), c_{x,y} + D_y(u)\} = \min\{2+5, 5+6\} = 7$$



4. **[Difficult]** Consider the network below. Suppose the link cost between x and y increases to 60. Recall that, in the lecture, we discussed that it will take a long time until the distance vector algorithm finally stabilizes ( $\geq 40$  iterations) due to the **count-to-infinity problem**. Describe a modification of the distance vector algorithm that correctly computes the distance vectors and avoids the count-to-infinity problem in this specific scenario.



# Recall: The Problematic Scenario


## link cost changes:

- node detects local link cost change
- “bad news travels slow” – **count-to-infinity problem:**

- y sees its adjacent link to x with new cost 60, but z has said it has a path to x at cost of 5. So y computes “my new cost to x will be 6, via z”; notifies z that it has a new cost of 6 to x.
- z learns that a path to x via y has new cost 6, so z computes “my new cost to x will be 7 via y”, notifies y of its new cost of 7 to x.
- y learns that a path to x via z has new cost 7, so y computes “my new cost to x will be 8 via y”, notifies z of new cost of 8 to x.
- z learns that path to x via y has new cost 8, so z computes “my new cost to x will be 9 via y”, notifies y of new cost of 9 to x.

...

➔ *Takes a long time to stabilize!*




Answer:

The technique to avoid the problem is called **poison reverse**:

### Key Idea:

- If z routes through y to get to destination x, then z will “lie” to y that its cost to x is infinity, i.e.,  $D_z(x) = \infty$  (even though z knows that  $D_z(x) = 5$ .)
- Since y believes that z has no path to x, it follows that y will never attempt to route to x via z, as long as z continues to route to x via y.




### How does it solve our problem?

- Suppose z has advertised  $D_z(x) = \infty$  to y.
- Time t0:
  - cost of (x,y) increases to 60.
  - y updates its vector and continues to route *directly* to x, although now at a higher cost of 60, i.e.  $D_y(x) = 60$ .
  - y informs z about its new DV.
- Time t1:
  - z receives y’s update. z updates its own DV and now uses the *direct* link to x, i.e.  $D_z(x) = 50$ .
  - z informs y about its new DV. **Note that z sends  $D_z(x) = 50$  and no longer  $D_z(x) = \infty$  !**
- Time t2:
  - y receives z’s update. y computes  $D_y(x) = 51$ . Now y will “lie” to z that  $D_y(x) = \infty$  because it routes through z. The algorithm stabilizes.

## 5. Describe how loops in paths can be detected in BGP.

Answer:

Since full AS path information is available from **an AS** to a **destination** in BGP, loop detection is simple – if a BGP peer receives a route that contains its own AS number in the AS path, then the use of that route would result in a loop.



- 6. Though we have link state and distance vector routing algorithms for least cost paths finding, the end-to-end path in today's Internet may not be the path with the best end-to-end performance. What are the possible reasons?

Answer:

**1. Intra-domain** routing protocols (i.e., iBGP) such as RIP and OSPF use the number of hops based or bandwidth-based cost metric, which does not necessarily correspond the end-to-end performance.

**2. Inter-domain** routing protocols (i.e., eBGP) are **policy driven**:

- For example, an ISP may route traffic along a less efficiency path to avoid paying an upper-tier ISP.
- Another example is that an ISP may not allow “transit” traffic to be routed through its network to avoid additional traffic overhead for which it does not generate any revenue.

